


Deep Learning and Artificial Intelligence WS 2024/25

Exercise 5: Recurrent Neural Networks

Exercise 5-1 Backpropagation through Time

Consider the following RNN:



Each state h_t is given by:

$$h_t = \sigma(W h_{t-1} + U x_t), \quad \sigma(z) = \frac{1}{1 + e^{-z}}$$

Let L be a loss function defined as the sum over the losses L_t at every time step until time T : $L = \sum_{t=0}^T L_t$, where L_t is a scalar loss depending on h_t .

In the following, we want to derive the gradient of this loss function with respect to the parameter W .

- (a) Given $\mathbf{y} = \sigma(W\mathbf{x})$ where $\mathbf{y} \in \mathbb{R}^n$, $\mathbf{x} \in \mathbb{R}^d$ and $W \in \mathbb{R}^{n \times d}$. Derive the Jacobian $\frac{\partial \mathbf{y}}{\partial \mathbf{x}}$!

[Solution: $\frac{\partial \mathbf{y}}{\partial \mathbf{x}} = \text{diag}(\sigma'(\cdot)) W$]

- (b) Derive the quantity $\frac{\partial L}{\partial W} = \sum_{t=0}^T \sum_{k=0}^t \frac{\partial L_t}{\partial h_t} \frac{\partial h_t}{\partial h_k} \frac{\partial h_k}{\partial W}$!

Exercise 5-2 Vanishing/Exploding Gradients in RNNs

In this exercise, we want to understand why RNNs are especially prone to the Vanishing/Exploding Gradients problem and what role the eigenvalues of the weight matrix play. Consider part b) of exercise 5-1 again.

- (a) Write down $\frac{\partial L}{\partial W}$ as expanded sum for $T = 3$. You should see that if we want to backpropagate through n timesteps, we have to multiply the matrix $\text{diag}(\sigma')W$ n times with itself.
- (b) Remember that any diagonalizable (square) matrix M can be represented by its eigendecomposition $M = Q\Lambda Q^{-1}$ where Q is a matrix whose i -th column corresponds to the i -th eigenvector of M and Λ is a diagonal matrix with the corresponding eigenvalues placed on the diagonals.¹

Proof by induction that for such a matrix the product $\prod_{i=1}^n M$ can be written as: $M^n = Q\Lambda^n Q^{-1}$!

¹Every eigenvector v_i satisfies the linear equation $Mv_i = \lambda_i v_i$ where $\lambda_i = \Lambda_{ii}$


(c) Consider the weight matrix $W = \begin{pmatrix} 0.58 & 0.24 \\ 0.24 & 0.72 \end{pmatrix}$. Its eigendecomposition is:

$$W = Q\Lambda Q^{-1} = \begin{pmatrix} 0.8 & 0.6 \\ -0.6 & 0.8 \end{pmatrix} \begin{pmatrix} 0.4 & 0 \\ 0 & 0.9 \end{pmatrix} \begin{pmatrix} 0.8 & -0.6 \\ 0.6 & 0.8 \end{pmatrix}.$$

Calculate W^{30} ! What do you observe? What happens in general if the absolute value of all eigenvalues of W is smaller than 1? What happens if the absolute value of any eigenvalue of W is larger than 1? What if all eigenvalues are 1?

Exercise 5-3 LSTMs

Recall the elements of a module in an LSTM and the corresponding computations, where \odot stands for pointwise multiplication.²



$$\begin{aligned} f_t &= \sigma(W_f h_{t-1} + U_f x_t) \\ i_t &= \sigma(W_i h_{t-1} + U_i x_t) \\ o_t &= \sigma(W_o h_{t-1} + U_o x_t) \\ \tilde{C}_t &= \tanh(W_c h_{t-1} + U_c x_t) \\ C_t &= f_t \odot C_{t-1} + i_t \odot \tilde{C}_t \\ h_t &= o_t \odot \tanh(C_t) \end{aligned}$$

- What do the gates f_t , i_t and o_t do?
- Which of the quantities next to the figure are always positive?

Let's now try to understand how this architecture approaches the vanishing gradients problem. To calculate the gradient $\frac{\partial L}{\partial \theta}$, where θ stands for the parameters (W_f, W_o, W_i, W_c), we now have to consider the cell state C_t instead of h_t . Like h_t in normal RNNs, C_t will also depend on the previous cell states C_{t-1}, \dots, C_0 , so we get a formula of the form:

$$\frac{\partial L}{\partial W} = \sum_{t=0}^T \sum_{k=1}^t \frac{\partial L}{\partial C_t} \frac{\partial C_t}{\partial C_k} \frac{\partial C_k}{\partial W}. \quad 3$$

- We know that $\frac{\partial C_t}{\partial C_k} = \prod_{i=k+1}^t \frac{\partial C_i}{\partial C_{i-1}}$. Let $f_t = 1$ and $i_t = 0$ such that $C_t = C_{t-1}$ for all t .

What is the gradient $\frac{\partial C_t}{\partial C_k}$ in this case?

²For a good explanation on LSTMs you can refer to <http://colah.github.io/posts/2015-08-Understanding-LSTMs/>

³The real formula is a bit more complicated since C_t also depends on f_t, i_t and \tilde{C}_t , which in turn all depend on W , but this can be neglected.