

Lecture Notes on
Deep Learning and Artificial Intelligence
Winter Semester 2024 /2025

Neural Networks and their
mathematical foundations

Script © 2024 Matthias Schubert



Chapter Overview

- Mathematical foundations:
 - vectors, matrices, and tensors
 - similarity, norms, and distance metrics
 - derivatives and gradients
 - 2nd order derivatives and extreme values
 - distributions and probability variables
- Basic Neural Networks
 - non-linearities
 - loss functions
 - weight initialization and input normalization
 - gradient descent
 - optimization techniques

Vectors, Matrices and Tensors

- Vectors, Matrices and Tensors $x \in \mathbb{R}^d, X \in \mathbb{R}^{n \times d}, \mathbb{X} \in \mathbb{R}^{d_1 \times \dots \times d_l}$

- inner product (dot product): $x^T u = \sum_{i=0}^d x_i \cdot u_i$

- outer product: $xu^T = \mathbf{x} \otimes \mathbf{u} = \begin{bmatrix} x_1 u_1 & x_1 u_2 & \cdots & x_1 u_m \\ x_2 u_1 & x_2 u_2 & \cdots & x_2 u_m \\ \vdots & \vdots & \ddots & \vdots \\ x_n u_1 & x_n u_2 & \cdots & x_n u_m \end{bmatrix}$

- matrix multiplication:

$$\begin{bmatrix} a_{11} & a_{12} & \cdots & a_{1n} \\ a_{21} & a_{22} & \cdots & a_{2n} \\ \vdots & \vdots & \ddots & \vdots \\ a_{m1} & a_{m2} & \cdots & a_{mn} \end{bmatrix} \times \begin{bmatrix} b_{11} & b_{12} & \cdots & b_{1p} \\ b_{21} & b_{22} & \cdots & b_{2p} \\ \vdots & \vdots & \ddots & \vdots \\ b_{n1} & b_{n2} & \cdots & b_{np} \end{bmatrix} = \begin{bmatrix} c_{11} & c_{12} & \cdots & c_{1p} \\ c_{21} & c_{22} & \cdots & c_{2p} \\ \vdots & \vdots & \ddots & \vdots \\ c_{m1} & c_{m2} & \cdots & c_{mp} \end{bmatrix}$$

$$c_{ij} = a_{i1}b_{1j} + a_{i2}b_{2j} + \cdots + a_{in}b_{nj} = \sum_{k=1}^n a_{ik}b_{kj}$$

Uses of Matrix Algebra

- linear function:

$$f : \mathbb{R}^d \rightarrow \mathbb{R}^n, f(x) = x^T W + b \text{ where } W \in \mathbb{R}^{d \times n}, b \in \mathbb{R}.$$

- cosine of angle α between two vectors $x, y \in \mathbb{R}^n$

$$\cos(\alpha) = \frac{x^T y}{\|x\| \cdot \|y\|}$$

- kernel matrix of a data set $X \in \mathbb{R}^{n \times d}$

$$K = X \cdot X^T \text{ with } x_{i,j} = X_i^T x_j$$

Metrics and Scalar Products

- Euclidian scalar product:

$$x, y \in \mathbb{R}^d : x^T y = \sum_{i=0}^d x_i \cdot y_i = \langle x, y \rangle$$

– properties:

- $\langle x, y \rangle = \langle y, x \rangle$ commutative
- $\langle x + y, z \rangle = \langle x, z \rangle + \langle y, z \rangle$ distributive
 $\langle x, y + z \rangle = \langle x, y \rangle + \langle x, z \rangle$
- Euclidian Norm: $\sqrt{\langle x, x \rangle} = \|x\|_2$
- Euclidian Metric:

$$Dist(x, y) = \|x - y\| = \sqrt{\langle x, x \rangle + \langle y, y \rangle - 2\langle x, y \rangle}$$

Differential Calculus

- given $f : \mathbb{R} \rightarrow \mathbb{R}$, the derivative is defined as:

$$f'(x) = \frac{d}{dx} f(x) = \lim_{\Delta \rightarrow 0} \frac{f(x + \Delta) - f(x)}{\Delta}$$

- given $f : \mathbb{R}^d \rightarrow \mathbb{R}$, the partial derivate is defined as:

$$\frac{\partial f}{\partial x_i} = \lim_{\Delta \rightarrow 0} \frac{f(x_1, \dots, x_i + \Delta, \dots, x_d) - f(x_1, \dots, x_d)}{\Delta}$$

- Jacobian of $f : \mathbb{R}^n \rightarrow \mathbb{R}^m$

$$\mathbf{J} = \frac{d\mathbf{f}}{d\mathbf{x}} = \left[\frac{\partial \mathbf{f}}{\partial x_1} \cdots \frac{\partial \mathbf{f}}{\partial x_n} \right] = \begin{bmatrix} \frac{\partial f_1}{\partial x_1} & \cdots & \frac{\partial f_1}{\partial x_n} \\ \vdots & \ddots & \vdots \\ \frac{\partial f_m}{\partial x_1} & \cdots & \frac{\partial f_m}{\partial x_n} \end{bmatrix}$$

2nd order Derivatives

- for $f : \mathbb{R} \rightarrow \mathbb{R}$: $\frac{\partial^2 f}{\partial x^2} = \frac{\partial f'(x)}{x}$

- for $f : \mathbb{R}^d \rightarrow \mathbb{R}$: $(\text{Hess } f)_{ij} = \frac{\partial^2 f}{\partial x_i \partial x_j} = \begin{bmatrix} \frac{\partial^2 f}{\partial^2 x_1} & \cdots & \frac{\partial^2 f}{\partial x_1 \partial x_d} \\ \vdots & \ddots & \vdots \\ \frac{\partial^2 f}{\partial x_d \partial x_1} & \cdots & \frac{\partial^2 f}{\partial^2 x_d} \end{bmatrix}$

- for $f : \mathbb{R}^n \rightarrow \mathbb{R}^m$: we a matrix per ouput channel

Optimization

- an optimization problem is defined as :

$$\operatorname{argmin}_{\theta} J(\theta)$$

where $\theta \in \mathbb{R}^p$ and $J : \mathbb{R}^p \rightarrow \mathbb{R}$

- **constrained optimization:**

$$\operatorname{argmin}_{\theta} J(\theta)$$

subject to

$$c_i(\theta) = 0, i \in \mathbb{I} \tag{1}$$

$$c_k(\theta) \geq 0, k \in \mathbb{K} \tag{2}$$

where $\theta \in \mathbb{R}^p$ and $J : \mathbb{R}^p \rightarrow \mathbb{R}$

Optimization and Gradients

- Convex functions:


for $0 \leq t \leq 1$ and $x_1, x_2 \in \mathbb{X}$:

$$f(tx_1 + (1 - t)x_2) \leq tf(x_1) + (1 - t)f(x_2)$$


- alternative definition, the Hessian is positive semidefinite
- convex functions have a unique minimum
- $J'(\theta) = 0$ yields the solution for $\operatorname{argmin}_{\theta} J(\theta)$

- Non-Convex functions:

- $J'(\theta) = 0$ could be a saddle point
- Second partial derivative test based on $\det(\operatorname{Hess} f(\theta))$ (<0 indicates saddle point)



https://en.wikipedia.org/wiki/Convex_function



https://en.wikipedia.org/wiki/Saddle_point

Finding the root of a function

2-order methods like Newton-Raphson (simple Newton)

- for univariate functions:

$$f : \mathbb{R} \rightarrow \mathbb{R} : f'(\theta_n) = \frac{f(\theta_n) - 0}{\theta_n - \theta_0} \Rightarrow \theta_0 = \theta_n - \frac{f(\theta_n)}{f'(\theta_n)}$$


- for multivariate functions:

$$f : \mathbb{R}^d \rightarrow \mathbb{R} : \theta_0 = \theta_n - \left(\frac{\partial f(\theta)}{\partial \theta \partial \theta^T} \right)^{-1} \frac{\partial f(\theta)}{\partial \theta}$$

Caution: For optimization, $f(\theta) = J'(\theta)$.

We need to compute the Hessian of $J'(\theta)$

to apply 2nd order optimization!



First order optimization

gradient descent or linear search:

- does not try to find the root of the derivative
- instead: walk along the gradient in small steps until it is 0.
- the derivative describes the slope of the tangent
- moving along the negative derivative for a small enough step α , leads to a smaller function value: $\theta_i = \theta_{i-1} - \alpha \cdot f'(\theta_{i-1})$

$$\exists \alpha \in \mathbb{R} : f(\theta_i) \leq f(\theta_{i-1})$$

If α is too big, I won't work properly

More, when we talk about training Neural Networks.

0-Order Optimization


- 1st and 2nd order optimizers require differential functions
- works on any function
- Idea: Instead of computing a gradient, sample k update steps and take the best ones. (Monte Carlo Search)

$$\theta_i = \theta_{i-1} - u \text{ where } u = \operatorname{argmin}_{u \in U} f(\theta_{i-1} - u)$$

- u is drawn from a random distribution
- the larger U the more likely we find an improvement.
- more sophisticated methods use populations of parameters and more fancy methods to generate updates
- for example: *Genetic Algorithms*

Probability Distributions

- sample space Ω
- for discrete Ω :
 - $\Pr(\omega)$ corresponds to the relative frequency when drawing ω from Ω .
 - $\sum_{i=0}^{|\Omega|} \Pr(\omega_i) = 1$ and $0 \leq \Pr(\omega_i) \leq 1$
- for continuous Ω :
 - $\Pr(\omega_i) = 0$ because $|\Omega| = \infty$
 - probability density function $p(\omega_i)$:
$$\int_{-\infty}^{\infty} p(\omega) d\omega = 1$$
 - cumulated Probability:
$$\Pr(\omega < x) = \int_{-\infty}^x p(\omega) d\omega$$



Mean and Variance

- a random variable X is a variable following a distribution over a sample space Ω .

- the expectation value of random variable X is:

$$\sum_{i=0}^{|\Omega|} p(x_i) x_i \text{ or } \int_{-\infty}^{\infty} p(x) x dx$$

- the variance of X is: $Var(X) = E \left((X - E(x))^2 \right) = \frac{\sum_{i=0}^{|\Omega|} (x - E(x))^2}{|\Omega|}$

- covariance of two random variable X, Y :

$$Cov(X, Y) = \frac{\sum_{i=0}^{|\Omega|} (x - E(X))(y - E(y))}{|\Omega|}$$

Mean and Variance for Random Vectors

- a random vector X consist of random variables x_i in each dimension $1 \leq i \leq d$

- mean vector: $E(X) = \begin{pmatrix} E(x_1) \\ \vdots \\ E(x_d) \end{pmatrix}$

- Co-Variance matrix of X :

$$Cov(X) = \begin{pmatrix} Var(x_1) & \dots & Cov(x_1, x_d) \\ \vdots & \ddots & \vdots \\ Cov(x_d, x_1) & \dots & Var(x_d) \end{pmatrix}$$

Estimating mean and covariance

- given a sample space Ω and data X drawn from the distribution $\Pr(x)$ over Ω :

- the empirical mean is computed as: $\hat{E}(X) = \frac{\sum_{i=0}^{|X|} x_i}{|X|}$

- sample bias: $\hat{E}(X) - E_{\Pr(x)}$

- covariance matrix of data matrix $X \in \mathbb{R}^{n \times d}$:

$$\frac{1}{|X|} (X - E(X))^T (X - E(X))$$

(normalized matrix product of the centered data matrix)

Estimating function parameters

given:

- observation $X \in \mathbb{R}^{n \times d}$, and labels $Y \in \mathbb{R}^{n \times l}$
- a prediction function $f_{\theta}(x) = \hat{y}$
- Likelihood function $\mathcal{L}(\theta, Y)$ describes how likely Y would be observed for parameters θ when feeding X to $f_{\theta}(x)$.

A maximum likelihood estimator determines θ^*

$$\theta^* = \operatorname{argmax}_{\theta} \mathcal{L}(\theta, Y)$$

Universal function approximators

- **Input domain:** $x \in X \subseteq \mathbb{R}^{d_1 \times \dots \times d_l}$
($l=1$ for vectors, $l=3$ for images, i.e., width, height, RGB channels)
- **Output domain:** $y \in Y \subseteq \mathbb{R}^{d_1 \times \dots \times d_l}$
(class probability, continuous prediction targets, images, text tokens, ..)
- A neural network is a parametric function $f_\theta: X \rightarrow Y$
where $\theta \in \Theta \subseteq \mathbb{R}^{d_1 \times \dots \times d_l}$ is called weights/parameters
- For a dataset $D = \{x_i, y_i\}_{i=1}^n \in X \times Y$ and $f(x_i, \theta) = \hat{y}_i$, we want to optimize an objective function $J_\theta: Y \times Y \rightarrow \mathbb{R}$ describing the similarity between \hat{y}_i and y_i for $1 \leq i \leq n$
 - as D is given, optimization needs to modify the weights θ :
for training, we consider $J_D(\theta)$ and in particular, its gradient $\nabla_\theta J_D(\theta)$

Neural Network Architecture

- the exact design of $f_{\theta}: X \rightarrow Y$ is called network architecture
- for now, we assume linear or dense functions:

$$f_{\theta}(x) = W^T x + b$$

- most architectures stack layers of functions:

$$f_{\theta}(x) = f^1(f^2(\dots f^n(x) \dots))$$

- stacking multiple linear layer functions results in a single linear function

⇒ separate layers by non-linear functions σ

$$f_{\theta}(x) = f^1(\sigma_2 f^2(\dots \sigma_n f^n(x) \dots))$$

σ is also called **non-linearity** or **activation function**

Non-Linearities



bounded logistic functions aka sigmoid functions:

- logistic function:

$$\sigma(x) = \frac{1}{1+e^{-x}}$$

- hyperbolic tangent:

$$\tanh(x) = 1 - \frac{2}{e^{2x}+1}$$



Rectified Linear Units (ReLU)

- ReLU:


$$\sigma(x) = \begin{cases} x & \text{if } x > 0 \\ 0 & \text{if } x \leq 0 \end{cases}$$

- Softplus:

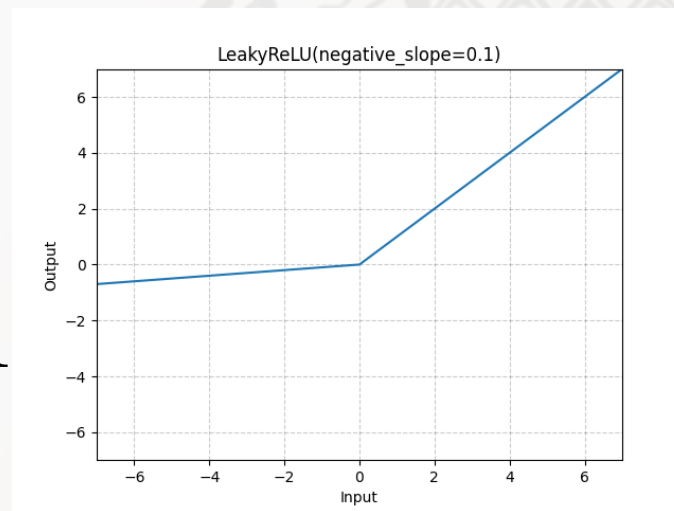
$$\sigma(x) = \frac{1}{\beta} \log(1 + e^{\beta x})$$

- leaky ReLU:

$$\sigma(x) = \begin{cases} x & \text{if } x > 0 \\ cx & \text{if } x \leq 0 \end{cases} \text{ with } c < 1$$



1.




2.

1. <https://pat.chormai.org/blog/2020-relu-softplus#>

2. <https://pytorch.org/docs/stable/generated/torch.nn.LeakyReLU.html>

Convex and Non-Convex Objective Functions

- $\forall (\theta_1, \theta_2) \in \Theta \times \Theta, \alpha \in [0,1] : J_D(\theta)$ is convex if
$$J_D(\alpha\theta_1 + (1 - \alpha)\theta_2) \leq \alpha J_D(\theta_1) + (1 - \alpha)J_D(\theta_2)$$



- for a convex function, any local minimum is global
- a univariate function is convex if the second derivative is non-negative for any parameter value
- a multivariate function is convex if the Hessian is positive semidefinite
(remember: the Hessian is the matrix of 2nd derivatives $\frac{\partial^2 f}{\partial x_1 \partial x_2}$)


Remarks on Convex Objectives

- objective functions for neural networks are usually non-convex
 - ⇒ convexity of $J_D(\theta)$ often depends on the complexity of $f_\theta(x)$
 - ⇒ optimisation of $J_D(\theta)$ often leads to local minima only
- optimal parameters θ^* are depending on D
 - ⇒ if you resample D , θ^* might not be optimal anymore.
- convexity can be exploited in optimisation algorithms

Objectives Continuous Outputs


for $y \in \mathbb{R}^n$ (if $y \in \mathbb{R}^{d_1 \times \dots \times d_l}$, we can flatten it to $\hat{y} \in \mathbb{R}^{d_1 + \dots + d_k}$)

- L1-loss: $L(y, \hat{y}) = \frac{1}{n} \sum_{i=1}^n |\hat{y}_i - y_i|$



- L2-Loss: $L(y, \hat{y}) = \frac{1}{n} \sum_{i=1}^n (\hat{y}_i - y_i)^2$

- Huber-Loss: $L(y, \hat{y}) = \frac{1}{n} \sum_{i=1}^n \begin{cases} \frac{1}{2} (\hat{y}_i - y_i)^2 & \text{if } |\hat{y}_i - y_i| < \delta \\ \delta \left(|\hat{y}_i - y_i| - \frac{1}{2} \delta \right) & \text{otherwise} \end{cases}$



Categorical Objectives (1)

Categorical data is usually encoded as a distribution over discrete outputs. For n categorical outputs, we consider $\hat{y} \in [0,1]^n$ and $y \in \{0,1\}^n$ as the output space:

Case 1: disjunctive categories (**Cross Entropy**)

- as y is a one hot encoding $\sum_{i=1}^n y_i = \sum_{i=1}^n \hat{y}_i = 1$.
- $f(x)$ is an arbitrary vector in \mathbb{R}^n
- use softmax function to resemble a quasi-one-hot distribution:

$$\hat{y}_i = \hat{f}_i(x) = \frac{e^{f_i(x)}}{\sum_{j=1}^n e^{f_j(x)}}$$

- Cross-Entropy Loss:

$$J_{\theta}(X, Y) = \mathbb{E}_{(x,y) \sim D} [-y^T \log(\hat{f}(x))] \approx -\frac{1}{k} \sum_{i=1}^k y_i^T \log(\hat{f}(x_i))$$

note: cross entropy only considers dimension j with $y_j = 1$.

If $\hat{f}_i(x) \rightarrow 0$ then $-\log(\hat{f}(x_i)) \rightarrow \infty$.

Categorical Objectives (2)

Case 2: overlapping categories (**Binary Cross Entropy**)

- each category needs to be treated independently $\sum_{i=1}^n y_i = \sum_{i=1}^n \hat{y}_i \leq n$
- each output dimension is scaled to $[0,1]$ by a logistic function

$$\hat{f}_i(x) = \frac{1}{1+e^{-f_i(x)}}$$


- Binary Cross-Entropy for k samples and n categories

$$\begin{aligned} J_{\theta}(X, Y) &= -\mathbb{E}_{(x,y) \sim D} [y^T \log(\hat{f}(x)) + ((1-y)^T \log(1 - \hat{f}(x)))] \\ &\approx -\frac{1}{k} \sum_{i=1}^k y_i^T \log(\hat{f}(x_i)) + (1 - y_i)^T \log(1 - \hat{f}(x_i)) \end{aligned}$$

note: $(1 - y) \log(1 - \hat{f}(x))$ is required to punish large values for categories x does not belong to ($y_j = 0$ but $\hat{f}(x_i) > 0$).

Multilayer Perceptron Architecture

neural network architecture consisting of n linear layers with an arbitrary non-linearity. (often 2 layers & logistic function)



$$f(x) = \sigma(W_h x^T + b_h) \quad \tilde{y}(x) = \sigma(W_{out} f(x)^T + b_{out})$$

Parameter Optimization

Various optimisation methods are available for different categories of objective functions:

- linear solvers (only for linear objectives)
- quadratic solvers (for quadratic objectives)
- Newton-Raphson (2. order optimisation requires the Hessian)
- ***gradient descent*** (1. order optimisation requires *gradient*)
- genetic and random search algorithms
(0. order optimisation: no requirements)
- various types of conjugate gradients
(various variants, good for sparse linear systems)

Note: Methods vary in resource requirements, convergence time, parameter sensitivity and whether they are guaranteed to converge.

Optimization for Neural Networks

- architectures are based on stacking non-linear functions
 - objectives are usually not convex
 - most architectures do not belong to a specific type of function for which dedicated optimisation algorithms are available
- architectures are generally continuously differentiable
 - methods like gradient descent and Newton-Raphson can be applied
 - Newton-Raphson requires the Hessian, which is quadratic in the number of parameters (infeasible for very large models)
 - if not, only genetic and randomised search is applicable

Gradient Descent (GD)

- recapitulate: the **derivative** measures how much a function $f(x)$ increases relative to a change in x , i.e., $f'(x) = \lim_{\Delta \rightarrow 0} \frac{f(x+\Delta) - f(x)}{\Delta}$ and the **gradient** is the derivative at a particular point x_0 .

- going into the opposite direction decreases $f(x)$:

$$\exists \alpha > 0: x_{t+1} \leftarrow x_t - \alpha f'(x_t) \Rightarrow f(x_{t+1}) \leq f(x_t)$$

- α is a hyperparameter called learning rate or step size
- adding the gradient is used for maximization (gradient ascent)
- gradient descent: iterate steps until $f(x_{t+1}) = f(x_t)$ or max. steps
- GD converges if $f(x)$ is differentiable and convex
- If $f(x)$ is not differentiable, we encounter an error when reaching an x where $f'(x_t)$ is undefined.
- If $f(x)$ is not convex, GD might not reach a global minimum but a local minimum or a saddle point.

Stochastic Gradient Descent (GD) Algorithms

Input: $D = \{x_i, y_i\}_{i=1}^n, \alpha, f(x, \theta)$

output: θ^* (local) optimal function parameter

init θ (*weight initialization*)

for e in episodes:

 for x_i, y_i in samples D:

$\hat{y}_i = f(x_i, \theta)$ #forward path

$L_\theta = J(\theta, \hat{y}_i, y_i)$ #compute loss

$\theta = \theta - \alpha \nabla_\theta L_\theta$ # backward path

return θ

GD for Neural Networks

- neural network $f(x, \theta)$ is a multi-layer non-linear function
- θ is a simplification for multiple groups of parameters

Example:

For $\tilde{y}(x) = \sigma(w_o(\sigma(W_i x^T + b_i))^T + b_o)$:

4 types of parameters: $\theta = (w_o, b_o, W_i, b_i)$

Each has a different derivative:

let $L_0^T = \sigma(W_i x^T + b_i)$

$$\frac{\partial \sigma(w_o L_0^T + b_o)}{\partial w_{o,i}} = \frac{\partial \sigma(w_o L_0^T + b_o)}{\partial w_o L_0^T + b_o} \frac{\partial w_o L_0^T + b_o}{\partial w_{o,i}} = \frac{\partial \sigma(w_o L_0^T + b_o)}{\partial w_o L_0^T + b_o} \frac{\partial \sum_{j=0}^n w_{o,j} L_j + b_o}{\partial w_{o,i}} =$$
$$\frac{\partial \sigma(w_o L_0^T + b_o)}{\partial w_o L_0^T + b_o} \left(\frac{\partial \sum_{j \neq i} w_j L_j + b_o}{\partial w_{o,i}} + \frac{\partial w_{o,i} L_i}{\partial w_{o,i}} \right) = \frac{\partial \sigma(w_o L_0^T + b_o)}{\partial w_o L_0^T + b_o} L_i$$

...

Computing Gradients in Neural Networks

- Neural Networks stack functions on top of each other
- use chain rule to compute gradients for parameter θ_i on the i -th layer:

- for 2 layers:
$$\frac{\partial f_1(f_0(x, \theta_0), \theta_1)}{\partial \theta_0} = \frac{\partial f_1(f_0(x, \theta_0), \theta_1)}{\partial f_0(x, \theta_0)} \frac{\partial f_0(x, \theta_0)}{\partial \theta_0}$$

- for d layers:


$$\frac{\partial f_d(f_{d-1} \dots f_0(x, \theta_0))}{\partial x} = \frac{\partial f_d(f_{d-1} \dots f_0(x, \theta_0))}{\partial f_{d-1} \dots f_0(x, \theta_0)} \frac{\partial f_{d-1}(f_{d-2} \dots f_0(x, \theta_0))}{\partial f_{d-2} \dots f_0(x, \theta_0)} \dots \frac{\partial f_0(x, \theta_0)}{\partial \theta_0}$$

note:

- for any θ_i the levels $i > 0$, $f_{i-1}(f_{i-2}(\dots), \theta_{i-1})$ yield a constant input
- intermediate results can be stored in the forward path and used for the gradients of the higher levels.
- gradients for higher levels are multiplied:
 - \Rightarrow if gradients at higher layers are all < 1 , the gradient might vanish
 - \Rightarrow if gradients at higher layers are all > 1 , the gradient might explode (this will become a problem in recurrent architectures)

Weight Initialization in Neural Networks

- multiplying gradients yields a problem for consistently small or large values
- the value of the gradient depends on the current parameters
- if parameters are not properly initialised, GD might start with unfavourable gradient sizes.



shaded blue areas
yield small gradients
 \Rightarrow slow convergence

Glorot Initialization

- weights are drawn iid from a distribution with 0 *mean*.
- non-linearity f is odd: $f(-x) = -f(x)$. (e.g. tanh)
- What is a good variance/range?
 - We want to achieve that the variance of signals z_i /gradients is the same between layers i where W, b represent parameter of a linear layer :

- $VAR[z_k^{i+1}] = Var[\sigma(W_{\circ,k}^i + b_k^i)]$ (forward)

- $VAR\left[\frac{\partial L}{\partial s_k^1}\right] = Var\left[\sum_{j=1}^{n^{i+2}} W_{j,k}^{i+1} \frac{\partial L}{\partial s_k^{j+1}} f'(s_k^i)\right]$ (backward)

- this can be achieved by: $VAR[W^i] = \frac{1}{n^i}$ (*forward*)

$$VAR[W^i] = \frac{1}{n^{i+1}} \text{ (backward)}$$

[Understanding the difficulty of training deep feedforward neural networks](#), Glorot et al. (2010)

Weight Initialization in Practice

odd non-linearities:

- for normal distribution:

$$W^i \sim N\left(0, \frac{2}{n^i + n^{i+1}}\right)$$

- for uniform distribution $U(-a, a)$:

$$a = \sqrt{\frac{6}{n^i + n^{i+1}}}$$

- Kaiming or He Initialization: (better for ReLU)


$$W^i \sim N\left(0, \frac{2}{\sqrt{n^i}}\right)$$

note:

- Frameworks like pytorch or tensorflow use these methods as default.
- Input x should be normalized between -1 and 1 to keep

Learning Rates and starting point

- learning rate α :
 - if α too big, $f(x_{t+1}) > f(x)$. GD might not converge.
 - if α too small, convergence is too slow to reach local minimum within given time frame.
- initialization x_0 :
 - might decide whether GD converges for non-convex functions.
 - the closer x_0 is to the minimum, the faster GD converges.
- 2nd order optimisation like Newton-Raphson choose the step size w.r.t. root of the derivative and avoid these problems. But computing the Hessian and its root is quadratic in the number of parameters.



<https://towardsdatascience.com/gradient-descent-algorithm-a-deep-dive-cf04e8115f21>

Variants of Gradient Descent

- **Stochastic Gradient Descent:** Compute loss on a single instance
- **Gradient Descent:** Compute Average Loss on the complete dataset D.
- **Batch Stochastic Gradient Descent:** Compute loss on a subset of D (batch) and average objective over this batch.
- gradient of the average is the average of the gradients:

$$\begin{aligned} \frac{\partial \frac{1}{n} \sum_{(x,y) \in D} J(x, y, \theta)}{\partial \theta} &= \frac{\partial \frac{1}{n} J(x_0, y_0, \theta) + \dots + \partial \frac{1}{n} J(x_n, y_n, \theta)}{\partial \theta} \\ &= \frac{1}{n} \frac{\partial J(x_0, y_0, \theta)}{\partial \theta} + \dots + \frac{1}{n} \frac{\partial J(x_n, y_n, \theta)}{\partial \theta} = \frac{1}{n} \sum_{i=0}^n \frac{\partial f(x_i, \theta)}{\partial \theta} \end{aligned}$$

Learning rates in Deep Learning

- optimal learning rate depends on the steepness of the derivative
- deep NNs have millions of parameters: derivatives are high-dimensional and steepness depends on the parameters
- improved optimizers for GD adjust the learning rate:
 - use individual learning rates for different parameters/layers
 - use momentum to stabilize the direction of the training (gradients are floating averages over the last couple of steps)
- most common optimizers today
 - AdaGrad
 - RMSPROP
 - ADAM

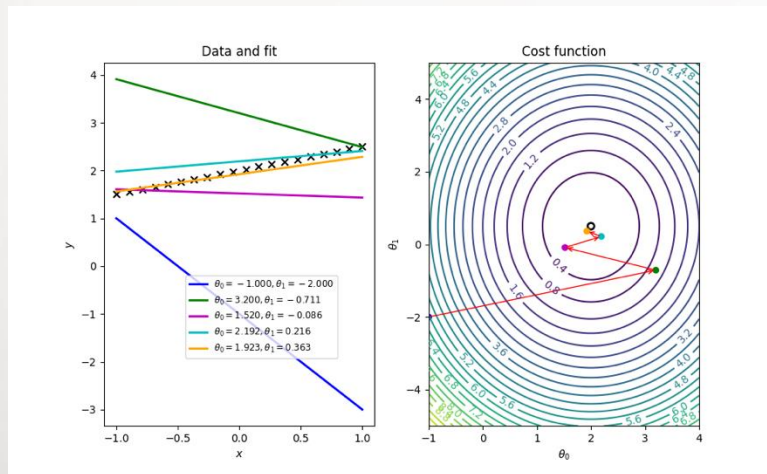
GD with Momentum

For too large learning rates α , GD starts to oscillate.

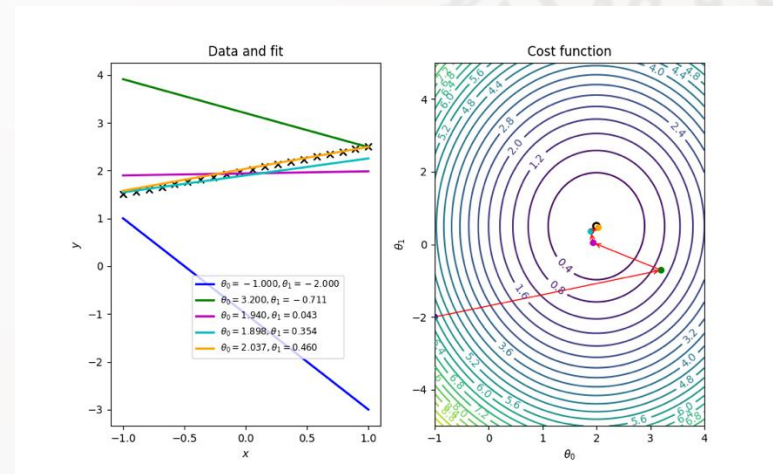
Idea: Update gradients with a floating average (momentum) to keep direction more stable. (γ : momentum, α : learning rate)

$$v_i^t = \gamma v_i^{t-1} + \alpha \frac{\partial f(\theta_i^t)}{\partial \theta_i^t},$$

$$\theta_i^{t+1} = \theta_i^t - v_i^t$$



$\gamma:0, \alpha:1.4$



$\gamma:0.1, \alpha:1.4$

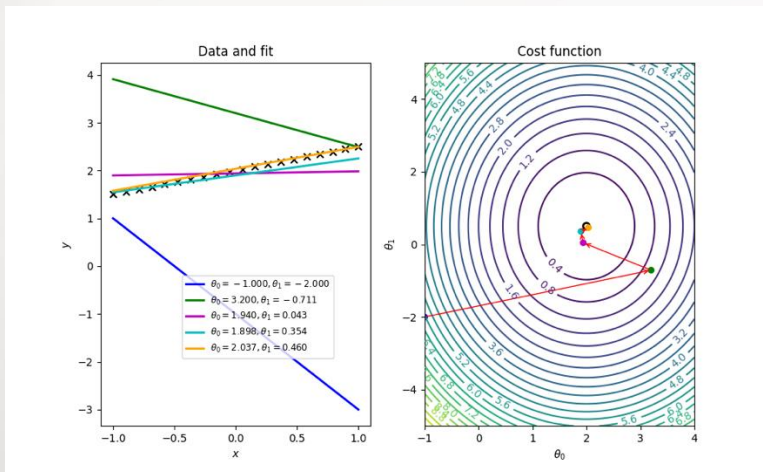
<https://distill.pub/2017/momentum/>

Gradient Descent with Nesterov Momentum

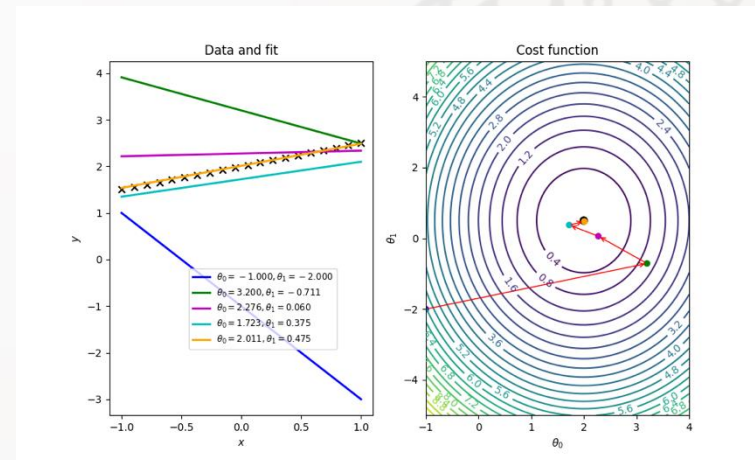
- gradient is computed after the update
- anticipates the update from momentum in the gradient

$$v_i^t = \gamma v_i^{t-1} + \alpha \frac{\partial f(\theta_i^t + \gamma v_i^{t-1})}{\partial \theta_i^t},$$

$$\theta_i^{t+1} = \theta_i^t - v_i^t$$



ordinary momentum $\gamma:0.1, \alpha:1.4$



Nesterov momentum $\gamma:0.04, \alpha:1.4$

AdaGrad

- introduces dedicated learning rates for each parameter
- learning rates are monotonically decreased
- computes the square sum of gradients for all parameters (large sum a lot of change happened \Rightarrow decrease LR)
- suitable for sparse data (no gradients of 0 inputs)

$$s_i^t = \sum_{\tau=1}^t \left(\frac{\partial f(\theta_i^\tau)}{\partial \theta_i^\tau} \right)^2$$
$$\theta_i^{t+1} = \theta_i^t - \frac{\alpha}{\sqrt{s_i^t + \epsilon}} \frac{\partial f(\theta_i^t)}{\partial \theta_i^t}$$

RMSProp

- individual learning rates with floating average

$$s_i^t = \gamma s_i^{t-1} + (1 - \gamma) \left(\frac{\partial f(\theta_i^t)}{\partial \theta_i^t} \right)^2$$

$$\theta_i^{t+1} = \theta_i^t - \frac{\alpha}{\sqrt{s_i^t + \epsilon}} \frac{\partial f(\theta_i^t)}{\partial \theta_i^t}$$

Adam (ADaptive Moment estimation)

- Uses momentum and RMSProp in combination
- currently, the gold standard

$$v_i^{(t)} = \beta_1 v_i^{(t-1)} + (1 - \beta_1) \frac{\partial f(\theta_i^{(t)})}{\partial \theta_i^{(t)}}$$

$$s_i^{(t)} = \beta_2 s_i^{(t-1)} + (1 - \beta_2) \left(\frac{\partial f(\theta_i^{(t)})}{\partial \theta_i^{(t)}} \right)^2$$

$$\theta_i^{(t+1)} = \theta_i^{(t)} - \frac{\alpha}{\sqrt{\frac{s_i^{(t)}}{1 - \beta_2^t} + \epsilon}} \frac{v_i^{(t)}}{1 - \beta_1^t}$$

note: the t-th power of β is used for bias correction

Newer Adam approaches

- **Adan[1]** extends Adam with Nestorov Momentum.
- **AdamW[2]** is a stochastic optimization method that modifies the typical implementation of weight decay in Adam, by decoupling weight decay from the gradient update.
- **Rectified Adam[3], or RAdam**, is a variant of the Adam stochastic optimizer that introduces a term to rectify the variance of the adaptive learning rate.

[1] <https://arxiv.org/abs/2208.06677v4>

[2] <https://arxiv.org/abs/1711.05101v3>

[3] <https://arxiv.org/abs/1908.03265v4>

Input Normalization

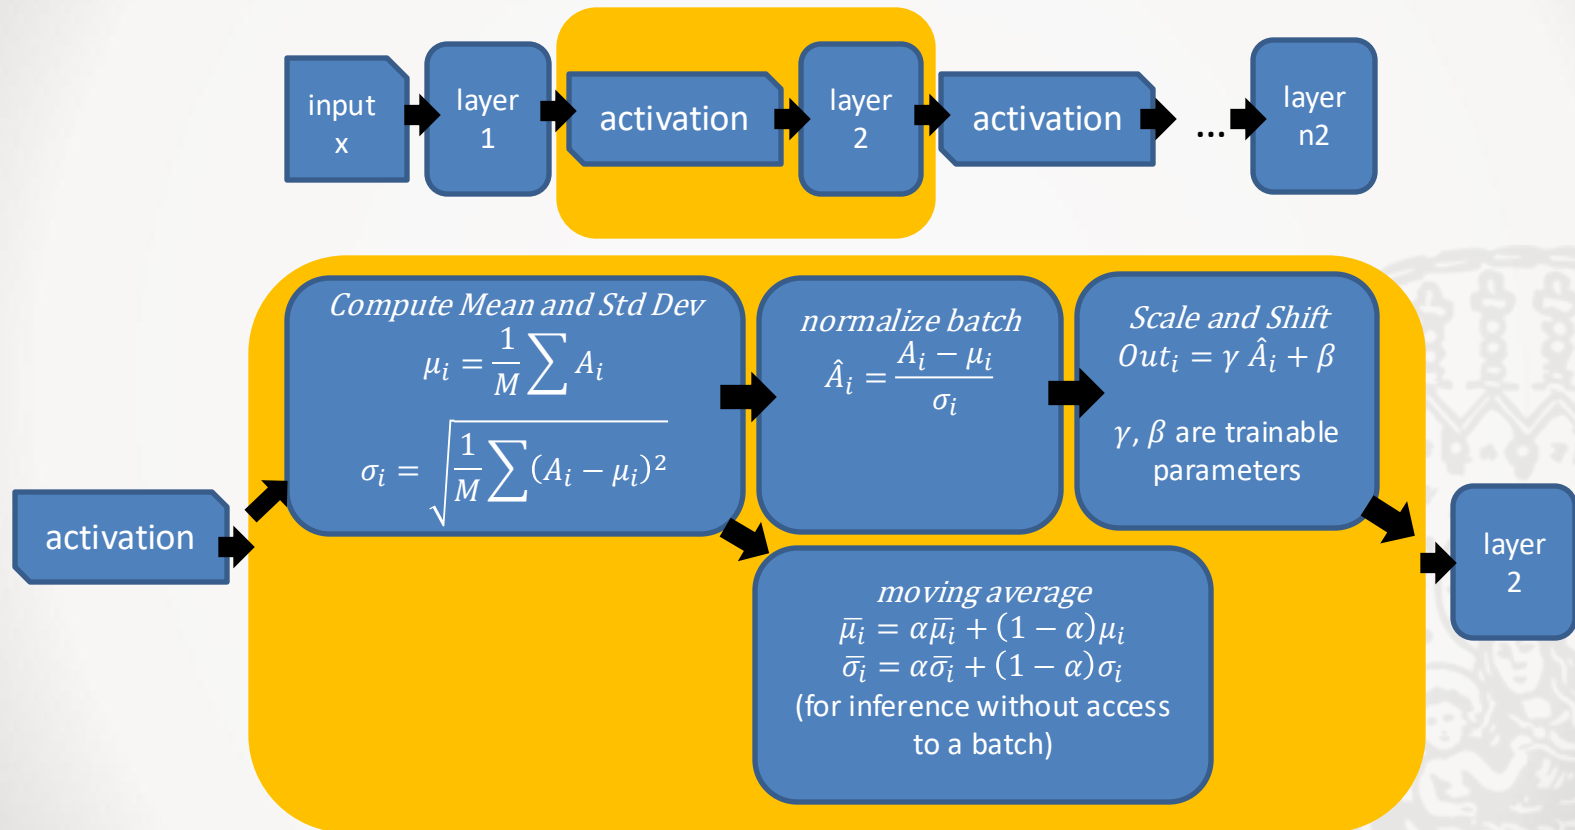
- inputs to the NN should be normalised to make all input features share the same scale (e.g. $\bar{x} = \frac{x-\mu}{\sigma}$)
- if raw features have varying scales, the gradients vary in scale as well. Updates progress at different speeds for varying parameters.
- normalisation counters this effect

Note:

- individual learning rates could be used here as well, but in Adagrad, RMSProp etc. this is not directly considered
- Input Normalisation helps with the input layer, but this effect could happen in internal layers as well

Batch-Normalization

batch normalisation layers adjust feature scaling for inner layers:



batch normalisation can be applied before or **after** activation.

Overfitting and Deep Architectures

- NN are universal function approximators with large amounts of parameters
⇒ model bias is pretty low
- generally, low-bias models allow to exactly fit the test data

So why do these models generalise to new data at all:

- SGD training helps:
 - gradient points to the optimum of the batch, not the complete data set, but the batch always changes
 - SGD takes a batch only for one step but does not converge to the minimum for one batch
- Usually, NNs are trained on very large data sets
- Deep architectures often contain layers for removing irrelevant information. (information bottlenecks, drop out, pooling etc.)

Note: Despite all, overfitting still might be a problem in NN.

Regularization via Dropout

What is a dropout layer?

- training: with probability $(1 - p_{keep})$ set an input value to zero \Rightarrow part of the network before dropout remains unused
- inference: all weights are used but with p_{keep} to keep the input scale

Why does it work?

- a way to reduce overfitting would be to train an ensemble of models, which is too expensive for larger architectures.
- usually, deep NNs yield a lot of parameters which implies that the network in itself encodes redundant information
- dropout emulates an ensemble by dynamically selecting subnetworks during training

Early Stopping

- SGD is a stepwise optimisation, and initial weights usually don't overfit but generalise well
 - in general, learning general concepts which hold for the majority of data decreases the joint loss stronger than fitting parameters to decrease the loss for particular instances
 - after the weights sufficiently represent the general concepts, fitting particular instances yields the majority of loss improvements
- ⇒ Early Stopping tries to end training at the sweet spot

Technically:

- extract a validation set to test performance independently
- if performance on the validation set consistently drops, stop training.
(depending on the batch size, several tries to improve might be allowed)

Summary

- building blocks: linear layers, non-linearities and loss functions
- gradients and optimization
- Gradient Descent, batch GD, Stochastic GD
- Momentum, Adagrad, RMSProp and ADAM
- Weight Initialization, Input Scaling and Batch Normalisation
- Overfitting, Dropout and Early Stopping